



# DynamoDB

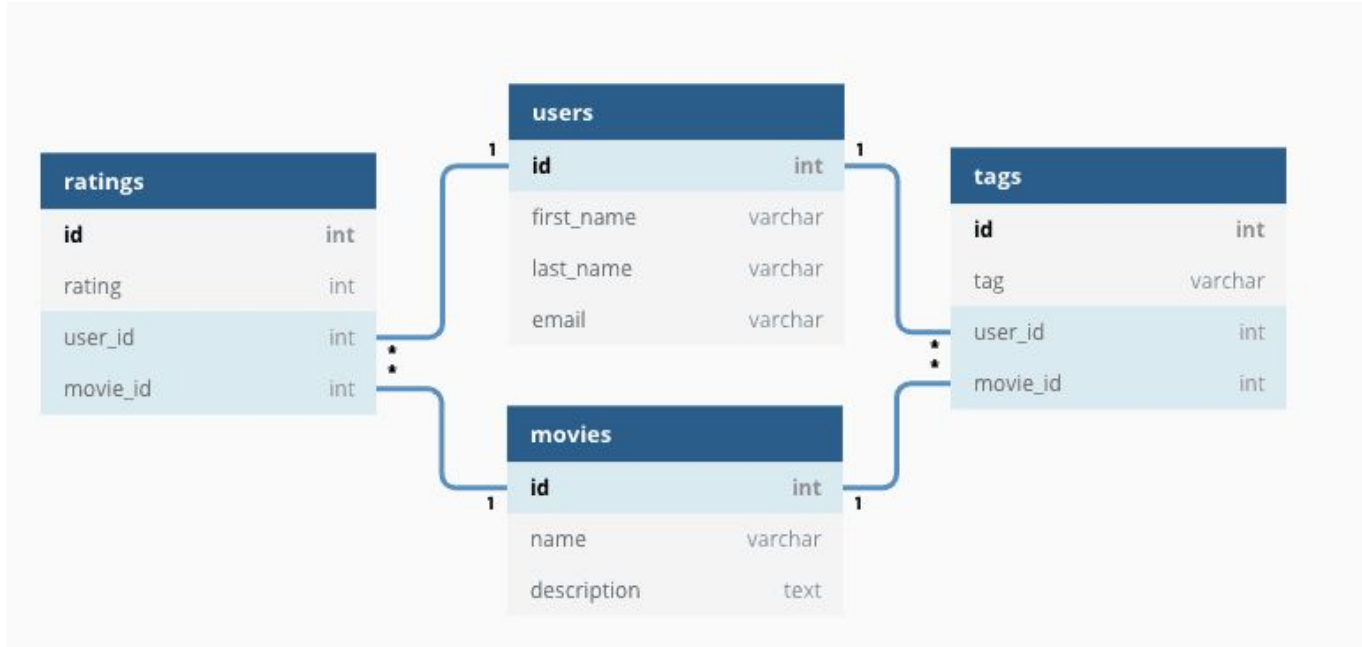
Part 1



# DynamoDB

Fast and flexible NoSQL database service for any  
scale

# Relational Databases




Relational Databases prioritized storage

# Non-Relational Databases



Non-relational Databases prioritize other aspects

# DynamoDB



Low Latency NoSQL Database  
It consists of Tables, Items and Attributes

# DynamoDB



Supports key-value and document data models  
DynamoDB is based on independent tables

# DynamoDB

## Two types of Keys

- Partition Key
- Sort Key




# Capacity Units



- **Read Capacity Units.-** One consistent read or two eventually consistent reads per second for items of up to 4Kb
- **Write Capacity Units.-** One write per second for items up to 1Kb




# Scan Operation



1. Returns all data in a table
2. Results can be filtered, but only once all items have been retrieved
3. Scans can use up the provision throughput for a large table in a single operation.

# Scan Operation



- 4. Retrieves data 1Mb at a time.
- 5. It can be configured to use Parallel Scan, that improves how long the operation takes

# Query Operation



1. Searches data in a table based on a Partition Key and an optional Search Key
2. By default eventually consistent



# DynamoDB

**DynamoDB Lab #1**

Query and Scan


# Local Secondary Index

- Can only be created when you are creating your table
- You cannot add, remove or modify it later
- Gives you a different view of your data based on the Sort Key




# Global Secondary Index

- Can be created with the main table or later on.
- Provides searching capabilities based on its Partition Key



# Indexes



- **Local Secondary Index.**- Same Partition Key different Sort Key
- **Global Secondary Index.**- Different Partition Key



# DynamoDB

## **DynamoDB Lab #2** Local and Global Indexes




# Provision Throughput

1. Your application needs to make 80 read requests per second
2. Each item is 3Kb in size
3. Consistent Reads are not needed

$$3\text{Kb}/4\text{Kb} = 0.75 = 1 \text{ RCU per request}$$
$$1\text{RCU} \times 80\text{requests/second} \times \frac{1}{2} = 40 \text{ RCU}$$


RCU = Read Capacity Unit,  
 $\frac{1}{2}$  because of eventually consistent reads

# DynamoDB On-Demand Capacity




1. DynamoDB scales your capacity units automatically based on the activity of your application
2. Great for unpredictable workloads

# DynamoDB Accelerator (DAX)



1. Fully managed, clustered, in-memory cache for DynamoDB
2. Delivers up to 10x performance only for Reads
3. Microsecond response times with million requests per second
4. Perfect for gaming or retail applications

# DynamoDB Accelerator (DAX)




- 5. DAX is a write-through cache implementation
- 6. DAX can help you reduce the read load in a DynamoDB table
- 7. It only works with eventually consistent reads

# Elastic Cache



1. In-memory Cache in the cloud
2. It sits between your database and your application
3. Can be used for databases/compute

# Types of Elastic Cache



**Redis**


- Snapshots and Replication
- Complex data structures
- Multi AZ Capability



**Memcached**


- Multi-Threaded
- No Multi AZ Capability

# DynamoDB Transactions



1. ACID Transactions (Atomic, Consistent, Isolated, Durable)
2. All or nothing operation


# DynamoDB TTL



1. Items are marked for deletion after a TTL
2. Once marked the item is deleted after 48h



# DynamoDB Streams



1. They record operations (insert, update, delete)
2. Before and After images are captured
3. Logs are created and encrypted up to 24h

# Provisioned Throughput Exceeded



1. When your request rate is too high
2. SDK will automatically retry requests until successful
3. You should use Exponential Backoff or reduce frequency

# DynamoDB Pricing

Pricing